Go, and a Simple, RAM-Based File System
Due: September 14, 2015, 11:59:59 pm

September 3, 2015

This project is designed to familiarize you with two technologies: the Go systems programming language from Google, and the FUSE interface for building user-level file systems. You will use Go and the Bazil/fuse language binding to build a simple user-space file system (UFS) that supports arbitrary file creations, reads, writes, directory creates, and deletions. The UFS will not support persistence, i.e. the contents of all UFS files will disappear the minute the UFS process is killed.

FUSE is a loadable kernel module; that is, a piece of code that is compiled and loaded directly into the kernel. Inside the kernel, potentially many different file systems co-exist, all living under the *virtual file system* layers, or VFS for short (see Figure 1). The VFS layer allows any application to call down into the file system with the same API, regardless of which, or which type of file system, is ultimately serving the file requests. We will use the kernel module, but *will never be hacking inside the kernel.*

The VFS layer knows which underlying file system to use based on which part of the file system namespace is being accessed. For instance, if you have an NFS volume mounted at /mnt/nfs, then the VFS layer knows to vector any calls to some file `/mnt/nfs/foo.c` to the NFS client, which will then communicate with a (possibly remote) NFS server for the actual data.

Likewise, we will mount a FUSE volume somewhere in the file system namespace (`/tmp/<userid>` is a good default), and any system calls to files under that directory get vectored to the FUSE module that lives under the VFS layer. In turn, the FUSE module will re-direct any calls sent to it back up to our user-level program. So all we have to write is the user-level program!

1   Setup

Ideally, you work on your own mac or linux box, as having root access makes things easier, and your mounts will not collide with anyone else's. However, I will be testing your code on `hyperion.cs.umd.edu` (Ubuntu), so you need to ensure that it works there. You will all have accounts on `hyperion` by Thursday 9/4, but not root access.

If you are working on your own machine, you might need to install FUSE yourself. This is definitely true if you are working on a Mac, where you must install osxfuse: http://osxfuse.github.io/. FUSE is part of the main Linux trunk for recent kernels, and so the kernel module is not needed for these kernels. More info at sourceforge.

Though it would not be too difficult, please do not attempt to get root, or to view or alter others' files on `hyperion`.

1.1   Installing Go

You will need to install two pieces of software: the Go language (often referred to as "golang", especially in web searches), and the Go language bindings for FUSE. The following is a good set of docs on both Go and bazil:

Figure 1: System Design

- http://golang.org

    - language spec
    - packages
    - the tour
    - effective go article
    - very good talk by Rob Pike
    - golang book

- Bazil/fuse

    - main site
    - talk

Install Go from here: http://golang.org/doc/install. The current version is 1.5. Best to install directly rather than use package installers such as `apt-get`, `macports`, or `homebrew`.

Once you've set up your directories and environment, you can install bazil/fuse:

```
go get bazil.org/fuse
```

This command will install bazil/fuse into your go path, irregardless of your current working directory.

1.2   Obtain the project files

You will retrieve all setup files from the class web page. Download the project tarball from: http://triffid.cs.umd.edu/dss/p1/p1.tgz under `$GOPATH/src`, then unpack with `tar xvfz p1.tgz`.

cd into `p1` and you will see two files: `hellofs.go` and `dfs.go`. The former is an example program from the bazil people that implements a static file system with a single file.

The latter is a skeleton of a dynamic in-memory file system. This code is actually runnable. Create a directory for a mountpoint, say `tmp/dss`. From the `p1` directory, type

```
go run dfs.go -debug /tmp/dss
```

Go into `/tmp/dss` in another shell window, create new subdirectories, and `cd` around. If this does not work, your environment is not set up correctly.

2   What to do

You will expand `dfs.go` into a fully functional in-memory file system. This means that files can be copied, created, edited, and deleted on mount points exported by your code. The `dfs.go` file has most of the boilerplate. You need to create objects that implement specific interfaces (`fs.Node` and `fs.FS`), pass those to fuse, and then fuse can call those objects' methods to implement file system functionality. You will do this through definition of the following methods, probably using less than 100 lines of code:

```
func (FS) Root() (n fs.Node, err error)
func (n *Node) Attr(ctx context.Context, attr *fuse.Attr) error
func (n *Node) Lookup(ctx context.Context, name string) (fs.Node, error)
func (n *Node) ReadDirAll(ctx context.Context) ([]fuse.Dirent, error)
func (n *Node) Getattr(ctx context.Context, req *fuse.GetattrRequest, resp *fuse.GetattrResponse) error
func (n *Node) Fsync(ctx context.Context, req *fuse.FsyncRequest) error
func (n *Node) Setattr(ctx context.Context, req *fuse.SetattrRequest, resp *fuse.SetattrResponse) error
func (p *Node) Mkdir(ctx context.Context, req *fuse.MkdirRequest) (fs.Node, error)
func (p *Node) Create(ctx context.Context, req *fuse.CreateRequest, resp *fuse.CreateResponse)
    (fs.Node, fs.Handle, error)
func (n *Node) ReadAll(ctx context.Context) ([]byte, error)
func (n *Node) Write(ctx context.Context, req *fuse.WriteRequest, resp *fuse.WriteResponse) error
func (n *Node) Flush(ctx context.Context, req *fuse.FlushRequest) error
func (n *Node) Remove(ctx context.Context, req *fuse.RemoveRequest) error
func (n *Node) Rename(ctx context.Context, req *fuse.RenameRequest, newDir fs.Node) error
```

Note that you do **not** have to use the boilerplate or data structures in `dfs.go`.

Note, however, that FUSE as a whole is documented extremely poorly, and the language bindings (implemented by third parties) are no exception. bazil/fuse is "documented" at http://godoc.org/bazil.org/fuse, but the code itself is probably a better resource. The best documentation will be in looking through the bazil/fuse code that will be calling your code. In particular, the following two files are most useful to understand:

```
$GOPATH/src/bazil.org/fuse/fuse.go
    and
$GOPATH/src/bazil.org/fuse/fs/serve.go
```

2.1   Running and Debugging

You can run Go programs a few different ways: I tend to just use the "go run" command above when debugging. To deconstruct this line, "go run" compiles/links/runs the program in `dfs.go`, "-debug" causes `dfs.go` to print lots of debugging information ("-fuse.debug" will cause the bazil/fuse shim layer to print out even more debugging information), and "/tmp/dss" is a directory I've pre-created to serve as a mountpoint.

A few last points:

- File systems are mounted onto directories in UNIX-like systems by *mount*-ing them. However, this is taken care of automatically when using the high-level Bazil interface `fs.Serve()`.

  Unmounting is another issue. Usually this is done like `sudo umount /tmp/dss`, but this requires sudo privileges. However, *dfs.go* includes a call to `fuse.Unmount()`, which appears to work on both my macs and a linux.
  Caveat: The unmount will not work if the old mount is still being used. For example, an unmount of `/tmp/dss` will fail if the working directory of any shell is inside the mount point.

  Run `dfs`, kill w/ Ctrl-c, and accessing the directory gives you a "Device not configured" error. However, running dfs again on the same mount point appears to correctly unmount the directory and re-use it. If this does not work, for some reason, the mountpoint will timeout and be useable after a few minutes anyway.

- You can add your own command-line arguments; see the implementation of the "-debug" flag in `dfs.go` as an example.

- You *may* attempt to debug with `gdb`, but `gdb` has poor support for high-level Go structures like channels, maps, and go-routines. That said, feel free to try: https://golang.org/doc/gdb. I would use this inside of `xemacs` or `Aquamacs`. Let me know if it works.

- Note that write calls might start at offsets *past* the current size of the file (`gcc` can do this, for example). This is valid; just initialize the space between the current size and the offset with zeros.

- Something wierd is happening with the `flush()` call in my implementation. Copying a file into a `dfs` mount causes `flush()` to be called. This is normal. However, if I then copy the file back out of the mount, the read causes another flush, and this is repeatable! Bonus points for explaining this behavior.

3   The details

You should have the following goals in developing your implementation:

1. allowing file/directory creation, manipulation, deletion.

2. allowing mounted files to be edited with emacs/vi (note: this should work w/ no problem as long as you don't delete the `Node.fsync()` method.).

3. allowing a large C project to be copied and built on top of your system.

Non-goals:

- file permissions

- links of either types

3.1   Deliverables, and grading

Code should be formatted as per Go conventions. If in doubt, run "go fmt" over your code.

I will grade each project out of 100 points, with up to 20 points coming from style. I do not care about comments per se, but I will give the most points to implementations that are the cleanest, simplest, and most efficient, probably in that order.

Submit your code via the submit server.

3.2   Timeliness

All of the projects have due dates. You will lose 10 points for every day your project is late. More importantly, note that the next project will be released the day the previous project is due.

3.3   Academic Integrity

You are all grad students. Please act like it.

You *may* discuss the project with other students, but you should not look at their code, nor share your own.

You *may* look at code and other resources online. However, if your code ends up looking like code from the web, and another student does the same, the situation is indistinguishable from cheating.

You *may* use the piazza website, and even post code. However, **you may not** post any code for the projects. `You may` post small snippets of generic Go illustrate or query some aspect of the language.

3.4 Zen

In general, I suggest the following approach for building projects in this class:

1. *keep algorithms simple* - no need to write a complicated search procedure when a simple linear search will suffice (or a sophisticated search procedure is callable from the standard library). We are interested in concepts and functionality in this class, not polish. We will be building proof-of-concept file systems, not commercial products.

2. *keep data representation simple*: Internal representations of files and directories is up to you. The skeleton in `dfs.go` takes an extremely simple approach using a single structure (the `Node`) to represent both. You could use distinct structures, but it seems to ramp up the complexity.